

LEVAV

MASTER PRODUCT REQUIREMENTS AND DUAL-AGENT EXECUTION SPECIFICATION

Version 4.1 | August 2026

Authoritative product, scientific, data, AI, engineering and delivery contract

Build edition: Claude Product Command + Codex Engineering Command

Confidential

CLAUDE: PRODUCT COMMAND	CODEX: ENGINEERING COMMAND
Owns product meaning, work packets, WRI and Levav 28 specifications, decisions, review and acceptance. Read-only over production code.	Owns all executable implementation, database changes, APIs, UI, tests, build, deployment and runtime debugging.

Core thesis: Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

Document Map

- Sections 1-7: authority, thesis, African problem, protected systems, product principles, language system and scope.
- Sections 8-12: identity, onboarding, Levav ID, Evidence Graph and WRI science.
- Section 13: Levav 28 adaptive workforce simulation and development environment.
- Sections 14-17: QuickWork, employed talent, Employer Intelligence and post-hire validation.
- Sections 18-30: opportunities, compensation, contribution, Learn, Champions, feed, Workforce Graph, data, monetisation, AI, trust and Africa-first infrastructure.
- Sections 31-39: current technology direction, data and API architecture, observability, metrics and product moat.
- Sections 40-50: Claude and Codex operating model, file ownership, handoffs, build sequence, Definition of Done and start instructions.
- Section 51: final product thesis.

1. Document Authority

This document is the single governing specification for the next stage of Levav.

It consolidates and supersedes the implementation authority of:

- Levav Final Master Product Requirements Document v3.0.
- Levav Master Product Requirements, System Architecture and Dual-Agent Execution Blueprint v3.3.
- Earlier product briefs, codebase audit notes and implementation assumptions where they conflict with this document.

The v3.0 product depth is preserved. The v3.3 execution discipline is preserved. The agent split is changed to match the strengths required for the build.

The authority order is:

1. This Master PRD v4.1 defines what Levav must become and the product rules that cannot be changed casually.
2. Approved Product Decision Records define later product decisions that explicitly amend this PRD.
3. The current repository and verified runtime define what exists today.
4. Implementation State documents report what is working, partial, mocked, broken or deferred.
5. Agent-specific instructions define how Claude and Codex operate. They do not override product requirements.

When documentation conflicts with running code, the agents must distinguish desired state from current state. They must not rewrite working systems simply because an older document described a different stack.

1.1 Mandatory audit vocabulary

Every existing feature or subsystem must be classified before major modification:

- KEEP: It already satisfies the requirement.
- ENHANCE: It works but needs stronger UX, data, intelligence, quality, security or performance.
- MODIFY: It exists but behaves differently from the approved requirement.
- COMPLETE: It is partially implemented.
- BUILD: It is missing.
- REMOVE: It conflicts with the approved product direction.
- DEFER: It is valid but outside the active release.

1.2 Non-negotiable preservation rule

Working authentication, organisation membership, database foundations, deployment configuration and security controls must be preserved unless a verified defect requires change.

No agent may replace a working architecture merely to match an example from an older PRD.

2. Executive Product Thesis

Levav is an African workforce intelligence platform that develops, measures and verifies workforce readiness.

Levav is not primarily a job board, social network, assessment website, learning marketplace or gig platform. Those surfaces exist to support one connected workforce intelligence system.

The governing principle is Person over Credential.

Qualifications remain useful. They are not enough. Levav must understand what a person can demonstrate, how that person behaves in real work contexts, how capability changes over time, what evidence supports each conclusion, and what happens after opportunity is given.

The product loop is:

Talent Identity -> Levav 28 -> Observed Behaviour -> Evidence Graph -> WRI -> Development -> Opportunity -> Verified Outcome -> Employer Intelligence -> Workforce Intelligence -> Better Levav 28 and WRI.

A feature is core only if it does at least one of these things:

- Develops readiness.
- Proves readiness.
- Creates or improves meaningful opportunity.
- Reduces employer uncertainty.
- Creates verified evidence.
- Improves workforce decisions.
- Learns from real outcomes.

3. The African Problem Levav Must Solve

3.1 Talent problem

Millions of young Africans enter the labour market while the supply of formal wage employment remains far smaller than the number of people seeking it.

The product must respond to two separate problems:

- A shortage of formal opportunities.
- A readiness and evidence gap that makes employers uncertain about who can perform well.

Levav must not teach a graduate to sit still until a perfect permanent vacancy appears.

It should help people remain economically and professionally active. A person should be able to develop capability, demonstrate it, offer productive capacity, build verified evidence and become more visible for the right opportunities.

3.2 Startup and SME problem

Many startups and growing organisations need work completed before they can afford full-time payroll.

They may need five hours of bookkeeping, two weeks of customer research, a three-month marketing coordinator, a fractional finance lead, a landing page, a reconciliation exercise, field data collection or operational support.

Levav must make this demand usable. QuickWork turns work that is too small, too early, too irregular or too specialised for a permanent hire into structured opportunity.

3.3 Employer problem

CVs and interviews can hide whether a candidate can think, prioritise, communicate, learn, collaborate and execute.

Levav replaces hiring by guesswork with evidence-first decision support.

3.4 Workforce institution problem

Traditional labour data can show participation, qualifications and employment status without showing demonstrated readiness, capability gaps, evidence quality or development trajectory.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

Levav should eventually provide privacy-protected aggregate workforce intelligence for governments, labour institutions, development organisations and research partners.

4. Product Promise by Stakeholder

Talent and professionals

Become ready. Prove readiness. Build evidence. Stay active. Become visible for suitable opportunities. Understand your development trajectory.

Startups and SMEs

Access verified productive capacity without taking unnecessary permanent payroll risk. Define work clearly. Find relevant people. Get work completed. Build a trusted history with talent.

Employers

Reduce hiring uncertainty. See the person beyond the CV. Understand Role Fit, WRI, evidence confidence and relevant outcomes. Make better interviews and better hires.

Employed professionals

Understand current-role readiness, next-role readiness, strengths, development gaps and trajectory without surrendering ownership of personal career data to an employer.

Governments and workforce institutions

Understand readiness, gaps and workforce patterns through governed aggregate intelligence rather than private individual surveillance.

5. Protected Core Systems

The following systems are protected product architecture. They should be built as connected parts of one ecosystem.

5.1 Personal Identity and Levav ID

One persistent person-owned identity. A person can activate multiple eligible workspaces without creating separate identities.

5.2 Levav 28

A personalised, adaptive workforce simulation and development environment that creates controlled evidence of demonstrated behaviour.

5.3 Evidence Graph

The factual backbone of Levav. Every meaningful readiness claim should be traceable to evidence, provenance, context, time and verification.

5.4 Workforce Readiness Index

A versioned interpretation of workforce readiness from evidence. WRI includes dimensions, confidence, evidence coverage and trajectory. It is not a personality score and is not a guarantee of employment success.

5.5 Levav Workforce Graph

A structured relationship model connecting people, roles, industries, skills, behaviours, evidence, learning, opportunities, organisations, compensation and outcomes.

5.6 QuickWork

The immediate capacity activation and verified real-world work engine.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

5.7 Employer Intelligence

The commercial decision-support layer for identifying and evaluating relevant talent and reducing hiring uncertainty.

5.8 Workforce Intelligence

The aggregate, privacy-protected intelligence layer for workforce planning and labour insight.

5.9 Levav Impact

A verified contribution and social-impact system connecting people, organisations, service, skills, outcomes and evidence without disguising unpaid labour as employment.

5.10 Professional Feed and Network

The professional activity and discovery layer for people, organisations, evidence, opportunities, learning, contribution and trusted workforce news. It must feel alive without turning popularity into readiness.

6. Product Principles

The following rules apply across every feature:

- Person over Credential.
- Evidence over claims.
- Development before ranking.
- Readiness is multi-dimensional.
- Score and confidence are separate.
- WRI and Role Fit are separate.
- Real-world evidence should increasingly validate or challenge simulated readiness.
- AI assists interpretation, simulation and workflow. AI does not become unreviewable truth.
- One person can have several legitimate work contexts without duplicate identities.
- Employer intelligence must be explainable and entitlement-controlled.
- Institutional analytics are aggregate by default.
- Paid visibility never overrides genuine relevance.
- Build for African work realities while maintaining global professional standards.
- Measure commitment without building surveillance.
- Low bandwidth is a first-class product condition.
- Mobile is a first-class product condition.
- Every protected score must be auditable internally.
- Every material AI decision path must be versioned and testable.

6.1 Product Language and UX Writing System

Levav must sound like one product from the public landing page to the deepest authenticated workflow. Language is part of the product system, not decoration added at the end.

LANG-001: One canonical product voice applies across every user-facing surface.

The voice must be clear, human, intelligent, confident and practical. It should respect African professional realities without using charity language, deficit framing, empty inspiration or exaggerated promises.

The same language standard applies to:

- Landing pages and campaign pages.
- Signup, onboarding and account recovery.
- Talent and employer dashboards.
- Levav 28 instructions, coaching and feedback.
- QuickWork briefs, assignments, reviews and disputes.

- WRI explanations, evidence states and confidence language.
- Feed, profiles, organisation pages and notifications.
- Levav Impact, Learn and Champions.
- Emails, SMS, alerts, empty states, error states and support copy.
- AI-generated summaries and assistants.

LANG-002: Canonical product terms must remain consistent.

Names such as Levav 28, QuickWork, Levav ID, Workforce Readiness Index, Evidence Confidence, Role Readiness, Role Fit, Evidence Graph, Levav Impact and Employer Intelligence must not drift into competing labels across different screens.

LANG-003: UX copy must be concise and action-oriented.

Acceptance criteria:

- Use sentence case for interface labels unless a brand name requires otherwise.
- Buttons describe the action the user is taking.
- Empty states explain what the user can do next.
- Error messages explain the problem and the next safe action.
- Scientific and AI outputs communicate uncertainty plainly and never overclaim.
- Avoid generic AI phrases, corporate filler and unnecessary jargon.

LANG-004: User-facing strings must be governable.

Codex should avoid scattering important product copy through components. Shared labels, product terminology and reusable UX messages should use an intentional copy architecture that can support future localisation and African languages.

LANG-005: Claude owns the language specification. Codex owns faithful implementation.

Claude may refine approved wording, terminology and voice rules. Codex must not silently rewrite product meaning while implementing screens.

7. Release Scope

7.1 Core release

The active product programme includes:

- Personal identity and Levav ID.
- Organisation membership and verification.
- Career taxonomy and intelligent onboarding.
- Professional profile and Personal Home.
- Levav 28.
- Evidence Graph.
- Versioned WRI.
- Employed-talent readiness journeys.
- QuickWork.
- Jobs and opportunity matching.
- Employer verification, search, subscriptions and hiring pipeline.
- Hire Confidence and post-hire Quality of Hire.
- Volunteer and Impact evidence.
- Levav Learn.
- Champions.
- Professional feed and activity layer.
- Compensation Intelligence foundations.
- Workforce Graph foundations.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Workforce Intelligence foundations.
- Admin, trust, moderation, privacy, audit and observability.

7.2 Deferred until justified

The following items must not distract the active build:

- A separate startup incubator product.
- A broad SkillsSpace blue-collar marketplace as a separate major surface.
- A broad public API marketplace.
- Deep enterprise SSO without a launch-client requirement.
- Unsupported cross-border salary coverage.
- Predictive labour forecasting before sufficient historical data exists.
- Specialised assessment engines outside the core WRI model before the baseline system is validated.
- Institutional data licensing before governance maturity.

8. Identity, Roles and Workspaces

AUTH-001: One individual has one persistent personal identity.

Acceptance criteria:

- Role switching cannot reset evidence, history or WRI.
- Talent capability is not represented by a duplicate account.
- Organisation access is membership-based.
- Organisation permissions are enforced server-side.
- A person may legitimately be talent, an organisation member, a client, a champion or an employed professional through related capabilities.

AUTH-002: Role and entitlement are separate concepts.

Acceptance criteria:

- Employer membership does not automatically unlock protected WRI.
- Subscription entitlement controls premium Employer Intelligence.
- Institutional access is separately contracted and scoped.
- Admin access is privileged and auditable.

AUTH-003: Personal evidence remains person-centred.

Acceptance criteria:

- Employment by an organisation does not grant that organisation blanket access to the person's private Levav data.
- Sponsored development uses explicit permission scopes.
- Evidence visibility is controlled independently from evidence existence.

9. Intelligent Onboarding and Personalisation

Onboarding must progressively understand the person. It must not feel like a long generic registration form.

The system should collect, confirm and progressively enrich:

- Goals.
- Current status.
- Career family.
- Current role.
- Target role.
- Seniority.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Industry.
- Work environment.
- Work history and CV.
- Skills, certifications, projects and references.
- Development goals.
- Opportunity preferences.
- Location and work mode.
- Compensation expectations where appropriate.
- QuickWork capacity where desired.

ONB-001: AI may infer profile structure, but the user confirms authoritative use.

Acceptance criteria:

- No silent career classification.
- Extracted CV fields are editable before confirmation.
- Inferences are visibly distinguishable from confirmed facts.
- Confirmed profile context informs simulations, matching and recommendations.

ONB-002: Career taxonomy must represent broad African economic sectors.

Acceptance criteria:

- The taxonomy supports technology, finance, accounting, agriculture, mining, healthcare, education, engineering, construction, hospitality, retail, sales, public administration, logistics, operations, creative work and other extensible sectors.
- Taxonomy values are admin-managed and versioned.
- Local job titles can map to canonical roles without erasing local language or context.

ONB-003: First-run value must appear early.

Acceptance criteria:

- The user can see an immediate next action after onboarding.
- A new user should understand why Levav 28 matters.
- A user interested in QuickWork should be able to declare capacity without completing irrelevant setup.
- Provisional states are clear when evidence is still weak.

10. Professional Profile and Levav ID

The profile is a professional evidence surface, not a decorated CV.

It may include:

- Professional photo and cover.
- Headline and current role.
- About.
- Experience and education.
- Skills.
- Projects and featured work.
- Certifications.
- Volunteer service.
- Verified employment evidence.
- Verified QuickWork history.
- Recommendations and structured reviews.
- Posts and professional activity.
- Selected badges and verified evidence.
- WRI visibility according to policy.

PROF-001: Claims and verified evidence must look different.**Acceptance criteria:**

- Verified items show source and verification status.
- Self-declared items are never visually presented as verified.
- Public profiles do not expose protected WRI by default.

PROF-002: Levav ID must become a durable professional identity anchor.**Acceptance criteria:**

- The identity remains stable through role changes.
- Evidence links back to the same person.
- Public sharing exposes only approved fields.
- Employer and institutional access follows policy rather than public URL availability.

PROF-003: A member can actively maintain a modern professional profile.**Acceptance criteria:**

- Edit profile photo and professional cover or background image.
- Edit headline, about, location, current role, experience, education, skills, projects, links and approved visibility settings.
- Feature selected work, verified evidence, Impact contributions and posts.
- Profile updates propagate consistently to the feed and other permitted surfaces.

PROF-004: Account identity settings support secure email change.**Acceptance criteria:**

- Email change requires re-authentication or equivalent secure verification.
- The new address must be verified before becoming authoritative.
- Existing evidence, WRI history, organisation membership and Levav ID remain attached to the same person.
- Email is an account credential and contact field, not a new identity.

11. Evidence Graph

The Evidence Graph is the factual backbone of Levav.

11.1 Evidence levels

E0 Self-Declared Evidence:

- CV claims.
- Profile statements.
- Declared skills.
- Unverified projects.

E1 Learning Evidence:

- Structured course completion.
- Learning assessments.
- Module outputs where measurement quality is sufficient.

E2 Controlled Simulation Evidence:

- Levav 28 scenario performance.
- Structured, versioned simulation tasks.
- Controlled behavioural observations.

E3 Verified Real-World Work Evidence:

- QuickWork completion.
- Verified projects.
- Verified employment outputs.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Structured volunteer placements.
- Verified supervisor reviews.

E4 Longitudinal Outcome Evidence:

- Probation outcomes.
- Retention.
- Rehire.
- Repeat engagement.
- Time-to-productivity.
- Performance reviews.
- Promotions.

EVD-001: Every evidence node records provenance.

Required fields include:

- Subject.
- Source.
- Evidence level.
- Evidence type.
- Verifier and relationship context where applicable.
- Time.
- Work context.
- Eligible WRI dimensions.
- Visibility.
- Status.
- Model or rubric version where generated through an instrument.

EVD-002: Evidence cannot be silently overwritten.

Acceptance criteria:

- Amendments preserve audit history.
- Disputes preserve the original record and resolution state.
- WRI recalculation is triggered through the controlled evidence pipeline.

EVD-003: Evidence quality is contextual.

Evidence strength may consider:

- Verification strength.
- Verifier trust.
- Recency.
- Duration.
- Complexity.
- Repeated pattern.
- Independence of sources.
- Context relevance.
- Measurement reliability.

One review must not dominate a long evidence history.

EVD-004: Factual observations and evaluations remain distinct.

Examples of factual fields:

- Check-in time.
- Submission time.
- Deadline compliance.
- Verified hours.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Completion state.
- Repeat engagement.

Examples of evaluation fields:

- Communication rating.
- Initiative rating.
- Quality judgement.
- Written feedback.

The WRI engine may weight these differently.

12. Workforce Readiness Index Scientific Methodology

WRI is a versioned interpretation layer over evidence.

It is not a personality score, intelligence score, social popularity score or guaranteed predictor of employment success.

12.1 Core dimensions

1. Critical Thinking: Analyses information, tests assumptions, evaluates alternatives and reaches reasoned conclusions.
2. Problem Solving: Defines problems, identifies causes, develops options, implements solutions and evaluates results.
3. Initiative and Ownership: Identifies useful work, acts appropriately, accepts responsibility and follows through.
4. Reliability and Execution: Fulfils commitments, manages deadlines, communicates delays and delivers.
5. Communication: Listens, asks useful questions and communicates clearly in relevant professional contexts.
6. Collaboration and Teamwork: Coordinates with others, manages disagreement and supports group outcomes.
7. Adaptability and Learning Agility: Responds to change, feedback and unfamiliar situations and improves over time.
8. Professional Discipline: Demonstrates preparation, time management, documentation, quality control and conduct.
9. Leadership Readiness: Demonstrates decision judgement, delegation, accountability, conflict handling and development of others at an appropriate level.
10. Contribution and Service Orientation: Creates verified value through contribution, support, service or mentorship beyond narrow self-interest.

12.2 Required WRI outputs

WRI-001: Store and present the following separately:

- Overall WRI estimate.
- Dimension profile.
- Evidence confidence.
- Evidence coverage.
- Readiness trajectory.
- Role Readiness where a target role exists.

WRI-002: WRI and Role Fit are separate.

A person may have strong general workforce readiness and weak fit for one role. Role Fit uses role requirements, relevant evidence, skills and Workforce Graph context. WRI describes demonstrated general readiness.

WRI-003: Confidence is separate from performance.

A high score with weak evidence must look different from a high score supported by diverse, recent evidence.

Confidence may increase while a score falls if stronger new evidence reveals a weakness.

Confidence inputs may include:

- Observation count.
- Evidence diversity.
- Independent sources.
- E3 and E4 coverage.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Recency.
- Duration.
- Cross-context agreement.
- Measurement quality.

WRI-004: Scoring coefficients remain configurable and versioned.

Acceptance criteria:

- Pilot weights are labelled provisional.
- No coefficient becomes permanent product truth without validation.
- Every snapshot stores model version, evidence window and timestamp.
- Historical snapshots never silently mutate after a model change.
- Role-specific weighting belongs in Role Readiness, not hidden changes to historical WRI.

WRI-005: Behaviourally Anchored Rating Scales are required where human or AI ratings are used.

Acceptance criteria:

- Anchors describe observable behaviour rather than vague adjectives.
- Rubric versions are recorded.
- AI evaluation uses structured outputs tied to the rubric.
- Low-confidence evaluation can trigger human review or a follow-up task.

12.3 Scientific validation programme

WRI must be treated as an evolving measurement system.

Required validation questions include:

- Content validity: Do tasks represent real work behaviour for the role and level?
- Construct validity: Does a dimension measure the intended behaviour rather than unrelated factors?
- Criterion validity: Does the measure relate to meaningful outcomes such as completion, rehire, probation, performance or retention?
- Test-retest reliability: Are estimates reasonably stable when underlying capability has not changed?
- Inter-rater reliability: Do qualified raters reach similar conclusions from the same evidence?
- Scenario equivalence: Are alternate scenarios comparable in difficulty and construct coverage?
- Fairness: Are outcomes systematically disadvantaging groups without job-related justification?
- Calibration: Do readiness levels correspond to observed outcomes over time?

12.4 Scientific governance

WRI model changes require governed approval.

The governance structure should include product leadership and qualified expertise in industrial-organisational psychology, psychometrics, data science, labour economics, HR, data protection and relevant sector domains.

WRI-GOV-001:

- Construct definitions are documented.
- Major scoring changes are versioned.
- Historical cohort tests run before release.
- Fairness analysis is part of model release review.
- Users have a dispute and review route for evidence that materially affects readiness conclusions.

13. Levav 28: Adaptive Workforce Simulation and Development Environment

Levav 28 must become one of Levav's defining product experiences.

It should feel like entering a workplace, not answering a questionnaire.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

The first 28-day journey establishes a meaningful baseline. Readiness continues evolving after Day 28.

13.1 Experience model

A user enters a persistent simulated work environment containing relevant elements such as:

- A manager.
- Colleagues.
- Clients or customers.
- Direct reports for leadership contexts.
- Projects.
- Deadlines.
- Messages.
- Documents.
- Data.
- Meetings or calendar pressure.
- New information.
- Scope changes.
- Feedback.
- Mistakes and recovery opportunities.

The simulated workplace should remember relevant journey context. Personas must behave consistently within the journey.

13.2 Personalisation contexts

Levav 28 must adapt to:

- Career family.
- Role.
- Industry.
- Seniority.
- Current status.
- Development goals.
- Prior evidence.
- Difficulty progression.

Required user contexts include:

- Graduate or unemployed professional.
- Employed professional.
- Senior management.
- Freelancer.
- Self-employed person.
- Student.
- Volunteer.

A Senior Accountant must not receive the same scenario as a graduate accountant. A senior leader should encounter delegation, strategy, conflict, ethics, resource allocation and executive communication contexts.

13.3 28-day architecture

Phase 1, Days 1-5: Baseline and Work Habits.

- Workplace onboarding and role context.
- Ambiguous brief and clarifying questions.
- In-basket prioritisation.
- Professional communication.
- Feedback intake and reflection.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

Phase 2, Days 6-10: Reasoning and Execution.

- Root cause analysis.
- Conflicting information.
- Resource allocation.
- Solution design.
- Logic and data interpretation.

Phase 3, Days 11-15: Complexity and Collaboration.

- Cross-functional coordination.
- Difficult stakeholder handling.
- Ethics and professional integrity.
- Delegation where seniority supports it.
- Evidence sufficiency checkpoint.

Phase 4, Days 16-21: Pressure and Adaptability.

- Scope change.
- Crisis communication.
- Mistake ownership and recovery.
- High-volume delivery under time pressure.
- Negotiation and trade-offs.
- Process improvement.

Phase 5, Days 22-27: Integrated Workplace Performance.

- Multi-day capstone work.
- Data analysis.
- Stakeholder alignment.
- Written proposal or work output.
- Execution planning.
- Presentation or defence where appropriate.
- Self-assessment against demonstrated evidence.

Phase 6, Day 28: Baseline Readiness Report and Next Plan.

- WRI dimension profile.
- Confidence.
- Evidence coverage.
- Strengths.
- Development priorities.
- Trajectory.
- Recommended next actions.
- Suitable QuickWork, jobs, Learn, volunteering or further simulations.

13.4 Simulation modalities

L28-001: The engine must support several work modalities rather than one chat pattern.

Target modalities include:

- Simulated inbox prioritisation.
- Manager and colleague chat.
- Client communication.
- Document review.
- Data table interpretation.
- Brief writing.
- Project planning.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Decision selection with reasoning.
- File or deliverable submission.
- Presentation or spoken response in later supported releases.

The release may phase modalities, but the domain model must not assume every scenario is plain chat.

13.5 Adaptive simulation

L28-002: Scenario difficulty and follow-up can adapt within governed limits.

Acceptance criteria:

- Adaptation is constrained by scenario schema and rubric coverage.
- A mid-task change can test adaptability without becoming arbitrary.
- The engine may introduce new stakeholder information, deadline changes, resource constraints or conflicting evidence.
- Scenario branches remain traceable to a published scenario version.
- Adaptation must not make two candidates incomparable without recording difficulty and scenario metadata.

13.6 Assessment mode and development mode

L28-003: Levav 28 blends assessment and development without confusing them.

Acceptance criteria:

- Not every activity affects WRI.
- Coaching tasks may give explicit guidance without producing scored evidence.
- Evidence-producing tasks declare eligible dimensions internally.
- Practice attempts are distinguishable from protected evidence attempts.
- Users can learn from feedback without being trapped by one weak performance forever.

13.7 Persistent workplace personas

L28-004:

- Manager, colleague, client and direct-report personas have role context.
- Persona state is isolated per user journey.
- Hidden rubrics never appear in user-facing prompts.
- Personas stay consistent with role, relationship, previous conversation and scenario facts.
- Persona responses may challenge the user without manufacturing facts that invalidate the assessment.

13.8 Evidence sufficiency checkpoint

L28-005: Day 15 opportunity expansion depends on evidence sufficiency, not calendar completion alone.

Acceptance criteria:

- The threshold is configurable.
- Click-through behaviour cannot unlock artificial readiness.
- Evidence quantity, construct coverage and quality are considered.
- A user below threshold receives targeted development tasks rather than a punitive dead end.

13.9 Day 28 report

L28-006:

- The report separates score, confidence and evidence coverage.
- It explains major strengths and development areas in plain language.
- It identifies what evidence is missing.
- It recommends a next path rather than presenting Day 28 as graduation from readiness development.
- The journey remains historically available.

13.10 Scenario authoring and quality control

L28-007: Levav requires an internal Scenario Studio or equivalent governed authoring workflow.

The system should eventually allow authorised product and subject-matter teams to:

- Define role and industry context.
- Define seniority.
- Select WRI dimensions.
- Define scenario facts.
- Define personas.
- Define difficulty.
- Define behavioural anchors.
- Define branch constraints.
- Define expected evidence artefacts.
- Run scenario QA.
- Version and publish.
- Retire or replace weak scenarios.

Generated scenarios are drafts until validated against the scenario schema and review process.

13.11 Integrity

L28-008: Protected evidence attempts must reduce easy answer injection without creating invasive surveillance.

Acceptance criteria:

- Paste and drag/drop can be restricted where appropriate.
- Timing and revisions may be recorded with consent and policy.
- Suspicious AI-like output is a risk signal only.
- No WRI penalty is based solely on an AI detector.
- Follow-up reasoning, alternate scenarios or review can resolve uncertainty.
- Network instability must not be confused with poor work discipline.

13.12 Low-bandwidth Levav 28

L28-009:

- Core text scenarios remain usable on weak connections.
- Non-proctored drafts can autosave locally and reconcile safely where appropriate.
- Media-heavy tasks have alternatives when the construct does not require media.
- Resumable state prevents loss of legitimate work.

14. QuickWork: Capacity Activation and Real-World Evidence

QuickWork is not a side-gig board.

It is Levav's immediate economic activation and E3 evidence engine.

It is designed for two realities:

- Talent may have useful capacity while waiting for or between formal roles.
- Startups and organisations may need capability before they can justify a permanent hire.

QuickWork should help build a culture of active productive capacity without devaluing professional identity.

14.1 Engagement types

QuickWork should support:

- Micro-task.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Hourly capacity.
- Fixed project.
- Short contract.
- Fractional role.
- Recurring support.
- Outcome-based assignment.

Examples can include finance, administration, customer support, operations, marketing, design, technology, research, data, sales support, project coordination and other verified categories.

14.2 Capacity Profile

QW-001: Talent can list exact productive capacity.

Acceptance criteria:

- Hours available today where useful.
- Hours per day, week or month.
- Specific days and time windows.
- Start and end dates.
- Remote, hybrid or on-site.
- Minimum and maximum engagement.
- Rates or pricing model.
- Work categories accepted.
- Location where relevant.
- Capacity can be paused or updated without deleting history.

14.3 AI Work Scoper

QW-002: QuickWork should help a client convert a vague need into a clear assignment.

The AI Work Scoper may assist with:

- Clarifying the desired outcome.
- Defining deliverables.
- Suggesting milestones.
- Identifying skills or role type.
- Estimating effort bands where supported.
- Identifying dependencies.
- Identifying what evidence can prove completion.
- Highlighting missing information.
- Producing a structured scope for client confirmation.

The AI may not invent a market price when valid rate data is unavailable.

14.4 Assignment lifecycle

QW-003: Core states are:

Draft -> Pending Verification -> Open -> Matched -> Funded or Confirmed -> Active -> Submitted -> Client Review -> Completed.

Exception states include:

- Disputed.
- Cancelled.
- Expired where appropriate.

Every transition must be authorised, timestamped and auditable.

14.5 Matching

QW-004: Matching must combine work need with real capacity.

Signals may include:

- Role and skill relevance.
- Relevant WRI dimensions.
- Evidence confidence.
- Similar verified work.
- Availability.
- Location and work mode.
- Rate compatibility.
- Reliability history.
- Client and talent eligibility.

Paid promotion must not override genuine fit.

14.6 Work behaviour evidence

QW-005: Capture factual work behaviour without surveillance.

Possible factual fields include:

- Agreed start time.
- Relevant check-in or session attendance.
- Hours committed and verified.
- Deadline and submission time.
- Completion status.
- Rework level.
- Delay communication.
- Repeat engagement.

Remote verification may use session attendance, assignment check-in, system activity or client confirmation.

Physical work may use client confirmation, QR check-in or consented location-assisted verification where appropriate.

Continuous invasive tracking is not required.

14.7 Structured closeout review

QW-006: Client review should measure relevant observable work behaviour.

Available review areas include:

- Punctuality.
- Reliability.
- Quality.
- Completion.
- Timeliness.
- Communication.
- Initiative.
- Problem solving.
- Critical thinking.
- Adaptability.
- Teamwork where relevant.
- Ownership.
- Professional conduct.
- Response to feedback.
- Resourcefulness.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Rehire willingness.
- Recommendation willingness.

The review form should adapt to the work type so clients are not forced to rate irrelevant dimensions.

14.8 Verified work artefacts

QW-007: Clients and talent may attach approved evidence.

Supported evidence may include:

- Photos of completed work.
- Screenshots.
- Documents.
- Reports.
- Designs.
- Approved outputs.
- Links.
- Before-and-after evidence where appropriate.

Every artefact has rights and visibility controls.

14.9 Commercial protection

QW-008: Payment and dispute rules must exist before commercial launch.

Required decisions include:

- Funding or commitment model.
- Escrow or equivalent protection where applicable.
- Milestones.
- Release conditions.
- Partial completion.
- Revision limits.
- Client non-response.
- Talent abandonment.
- Cancellation.
- Refund rules.
- Dispute workflow.
- Chargeback handling.
- Platform fees.
- Transaction states.

14.10 Client reputation

QW-009: Clients build reputation as well as talent.

Signals may include:

- Payment reliability.
- Review turnaround.
- Dispute rate.
- Repeat hires.
- Assignment clarity.
- Verified organisation status.

14.11 Opportunity progression

QW-010: QuickWork should support progression from small work into deeper opportunity without forcing it.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

Possible pathways include:

- Repeat assignment.
- Recurring capacity.
- Short contract.
- Fractional role.
- Formal hiring pipeline.

A successful QuickWork relationship may become evidence for a larger role decision.

15. Employed Talent Journey

Employed users must not be treated as entry-level job seekers.

EMP-TAL-001: Current Role Readiness answers how strongly current evidence supports effective performance at the user's present level.

EMP-TAL-002: Next Role Readiness answers what evidence supports a target role or career level.

EMP-TAL-003: Development Trajectory shows which dimensions are improving, stable or weakening.

EMP-TAL-004: Verified Employment Evidence may include role confirmation, supervisor verification, project completion, promotions, leadership responsibility, performance review or approved work output.

Acceptance criteria:

- Confidential employer information is not required to prove work.
- Employer verification requires user consent.
- The person's Levav identity remains person-owned.

16. Employer Intelligence

Employers do not primarily pay Levav for profile access.

They pay to reduce hiring uncertainty.

16.1 Employer onboarding

EMP-001:

1. A representative uses a personal identity.
2. The representative creates an organisation or requests membership.
3. Organisation details are verified through available authoritative sources or a manual process.
4. The employer defines industry, locations, size, hiring markets and goals.
5. The employer creates a Talent Search Profile.
6. Natural language requirements may be converted into structured criteria for confirmation.
7. A verified free employer sees a limited relevant preview with protected WRI locked.
8. Subscription unlocks entitled intelligence.

16.2 Talent Search Profile

EMP-002: Employer search begins with the person the organisation actually needs.

Search criteria may include:

- Role.
- Seniority.
- Industry.
- Location.
- Work mode.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Skills.
- Relevant experience.
- Availability.
- WRI dimensions where lawful and legitimate.
- Evidence confidence.

Employers may weight legitimate job-related dimensions such as critical thinking, initiative, reliability, communication, leadership or teamwork.

Protected characteristics are excluded from employment ranking.

16.3 Explainable matching

EMP-003: Every recommended candidate must have an explainable reason for appearing.

The employer should be able to see:

- Role Fit.
- Relevant skills and experience.
- Relevant WRI dimensions according to entitlement.
- Evidence confidence.
- Evidence gaps.
- Availability where applicable.

The system must not present one opaque magic score as the entire hiring decision.

16.4 Employer paywall

EMP-004: Access states include:

Public:

- Public professional information.
- No protected WRI.

Verified Employer Free:

- Organisation profile.
- Talent Search Profile.
- Limited relevant talent preview.
- WRI locked.

Subscribed Employer:

- Advanced search.
- Entitled WRI depth.
- Evidence context.
- Candidate comparison.
- Saved pools.
- Messaging and hiring intelligence.

Enterprise:

- Seats.
- Deeper analytics.
- Compensation Intelligence where data supports it.
- Quality of Hire.
- Workforce analytics.
- Contracted integrations.

Every protected WRI view requires server-side entitlement checks and an audit event.

16.5 AI Job Design Assistant

EMP-005: The assistant may analyse a proposed role before publication.

It can:

- Identify unclear or conflicting responsibilities.
- Convert free text into structured role requirements.
- Flag unrealistic combinations of seniority, experience, skills or salary only where valid evidence supports the observation.
- Suggest clearer outcomes and success measures.

The employer confirms changes.

16.6 Interview Intelligence

EMP-006: Interview guidance uses evidence gaps rather than generic questions.

Acceptance criteria:

- Highlight strong evidence.
- Highlight uncertain areas.
- Suggest role-relevant questions tied to evidence gaps.
- Never recommend automatic rejection solely because of WRI.

17. Hiring Intelligence and Post-Hire Validation

HIRE-001: Pre-hire output is Hire Confidence or Predicted Fit, not guaranteed ROI.

It may use:

- Role Fit.
- Relevant WRI dimensions.
- Evidence confidence.
- Evidence coverage.
- Trajectory.
- Relevant verified outcomes.

Major contributing factors and limitations must be explainable.

HIRE-002: Post-hire Quality of Hire uses employer-confirmed outcomes.

Possible inputs include:

- Probation outcome.
- Retention.
- Time-to-productivity.
- Performance.
- Manager feedback.
- Promotion or expanded responsibility later.

Predicted Fit can then be compared with observed outcome.

HIRE-003: Employer outcome participation should return value.

Employers who contribute lawful post-hire outcomes may receive:

- Calibration reports.
- Quality of Hire analytics.
- Better internal benchmarks.
- Plan-based incentives.

HIRE-004: Hiring ROI is calculated only when sufficient employer financial inputs exist.

The formula and assumptions are visible. The result is labelled as an estimate unless inputs are confirmed.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

18. Jobs and Opportunity Matching

Jobs remain an opportunity surface within the readiness system.

JOB-001:

- Jobs map to role taxonomy.
- Job requirements map to skills and legitimate WRI dimensions.
- Candidates can understand major matching reasons.
- Applications retain the evidence snapshot relevant at the time of decision where lawful.
- Employer pipeline state is auditable.

JOB-002: Opportunity matching should use evidence rather than keyword matching alone.

JOB-003: A job application must not erase QuickWork, Levav 28 or other verified evidence. The system should combine relevant signals.

19. Levav Compensation Intelligence

AI may interpret compensation data. It may not invent numerical benchmarks.

COMP-001: Numerical benchmarks require valid data such as:

- Levav-observed market data.
- Employer-reported data with governance.
- Licensed datasets.
- Valid public or partner datasets.

COMP-002: Every benchmark exposes context:

- Role mapping.
- Geography.
- Observation period.
- Coverage indicator where available.
- Data source category.
- Confidence.

COMP-003: Insufficient data is an acceptable product state.

The system may return:

- A broader benchmark.
- A low-confidence result.
- Insufficient data.

It must never fabricate a salary range.

Potential capabilities include:

- Market range.
- Median and percentile bands where valid.
- Salary bands.
- City or country comparison where supported.
- Evidence-based scarcity premiums.
- Offer guidance.
- Internal equity with authorised employer data.
- QuickWork rate intelligence.
- Simpler talent-facing market context.

20. Levav Impact: Verified Contribution and Social Impact

Levav Impact is a first-class contribution system. It allows people to use real capability in verified service contexts, allows organisations to mobilise skilled contributors, and turns legitimate contribution into a visible record of work, service and outcomes.

It must never turn volunteering into disguised unpaid employment. Contribution is distinct from employment, and the user must understand which pathway they are entering.

20.1 Impact opportunities

IMPACT-001: Verified organisations can publish structured Impact opportunities.

Supported opportunity types may include:

- Skilled volunteering.
- Community projects.
- Mentorship and professional service.
- Pro bono professional assignments.
- Time-bound campaigns and field initiatives.
- Organisational service placements with defined outputs.

Required fields include cause or purpose, skills, location or work mode, duration, expected outputs, supervisor, safety/context information and whether any costs or stipends are covered.

20.2 Verified contribution record

IMPACT-002: Completed service can create a verified contribution record.

Acceptance criteria:

- Organisation confirms role, dates, hours or duration, outputs and completion.
- Relevant artefacts may be attached under visibility and rights controls.
- The contributor can dispute inaccurate records.
- Verified Impact appears on the professional profile when the user permits it.
- Only behaviourally relevant, sufficiently verified observations may contribute to eligible WRI dimensions.
- Impact participation alone never creates an automatic WRI increase.

20.3 Organisation Impact workspace

IMPACT-003: Verified organisations can manage contribution programmes and outcomes.

The workspace may include opportunity creation, applicants or invited contributors, placement status, supervisor verification, outputs, completion, contribution hours, approved outcome measures and reporting.

20.4 Impact identity and discovery

IMPACT-004: Impact should strengthen professional identity without becoming a vanity leaderboard.

Users may feature causes, verified service, selected outcomes, testimonials and contribution stories. Organisations may publish verified Impact updates into the professional feed. Popularity metrics must not substitute for verified contribution.

20.5 Impact measurement guardrails

IMPACT-005: Levav distinguishes activity, output and outcome.

Hours served are activity. A completed deliverable is an output. A claimed change in a community or beneficiary group is an outcome only when the organisation has a legitimate measurement basis. The platform must not manufacture social-impact claims.

Impact success metrics may include verified placements, completion, repeat participation, outputs delivered, verified contribution hours, skills deployed and organisation-confirmed outcomes where measurement quality is sufficient.

21. Levav Learn

Learn exists to close evidence-backed development gaps.

LEARN-001: Courses may be free, paid, sponsored or employer-funded.

LEARN-002: Course completion alone does not automatically increase WRI.

Learning can contribute evidence only when the assessment evidence is valid for an eligible dimension.

LEARN-003: Recommendations explain why learning is relevant.

LEARN-004: Learning effectiveness should be measured where possible through pre/post evidence or relevant WRI changes, without claiming causality beyond the data.

22. Champions

Champions are verified high-value contributors, not a generic influencer badge.

CHAMP-001: Entry can occur through invitation, nomination or application.

CHAMP-002: Verification may consider:

- Professional credibility.
- Verified portfolio.
- Expertise.
- Learning contribution.
- Supported causes or community contribution.

Champion status must not inflate WRI or search ranking merely because the badge exists.

23. Professional Feed and Network

The feed must feel active, useful and professional. It is the daily discovery layer of Levav, connecting identity, work, evidence, learning, organisations, Impact, opportunities and relevant workforce information.

It is not a vanity social network and social popularity is never a readiness score.

23.1 Personal and organisation presence

FEED-001: Every eligible member can maintain a recognisable professional presence.

Acceptance criteria:

- Profile photo and professional cover or background image.
- Headline, current role, location, about and selected expertise.
- Featured work, selected evidence, Impact contributions and pinned or featured posts.
- Edit-profile controls connect to the canonical professional profile rather than creating a duplicate social identity.
- Verified organisations have organisation pages and can publish updates under controlled membership permissions.

23.2 Publishing

FEED-002: Members and authorised organisations can create useful professional posts.

Supported post types should include:

- Text updates.
- Images and compressed media.
- Project and portfolio updates.
- Documents or approved links where safe.
- Verified QuickWork completion or work evidence cards where visibility permits.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Levav 28 milestones that reveal no protected rubrics or scores.
- Learn achievements and useful professional knowledge.
- Levav Impact stories and verified contribution updates.
- Employer, organisation and opportunity announcements.

Users should be able to edit or delete their own posts subject to audit and evidence-preservation rules where a post has been used as a verified artefact.

23.3 Engagement and network

FEED-003: Support follow, relevant reactions, comments, replies, save, share or repost where appropriate, reporting and notification controls.

Mentions and topic tags may support discovery. Anti-spam, blocking, muting and reporting controls are required before broad public growth.

23.4 Feed discovery

FEED-004: The feed should provide useful discovery without rewarding noise.

Candidate feed views may include Following, For You, Opportunities, Learning, Impact and organisation updates. Personalisation may use declared interests, career context, follows, location, skills and permitted activity signals.

WRI, protected assessment data and hidden employer intelligence must not be exposed or used as public social-ranking signals.

23.5 News and What is New

FEED-005: Levav can surface a trusted What is New layer.

This may include Levav product updates, employer announcements, verified organisation news, workforce and career insights, and curated external professional news where sourcing is clear. External news must retain source attribution and link to the original publisher. AI summaries must not invent facts.

23.6 Trust and verified context

FEED-006: Verified context should be visible where useful.

A post may display that a work item, organisation, achievement or Impact contribution is verified, but the interface must distinguish a verified source from a claim that the entire post is objectively true.

23.7 Feed quality and low-bandwidth behaviour

FEED-007: The feed is mobile-first, fast and usable on constrained networks.

Acceptance criteria:

- Lazy-load media and prioritise text.
- Compress images and use responsive sizes.
- Provide low-data behaviour and graceful media failure.
- Use pagination or controlled infinite loading that does not trap the user or consume unnecessary data.
- Preserve compose drafts through ordinary connection interruptions where practical.

23.8 Feed integrity

FEED-008: Popularity is not readiness.

Likes, reactions, follower counts, posting frequency and viral reach must never directly increase WRI, Role Fit or employer ranking. Sponsored content, if introduced, must be clearly labelled and cannot override employment relevance.

24. Levav Workforce Graph

AI should operate over structured workforce context rather than inventing Levav's ontology dynamically.

Core node families include:

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Person.
- Role.
- Skill.
- Behaviour or WRI dimension.
- Organisation.
- Opportunity.
- Learning.
- Evidence.
- Outcome.
- Compensation context.

Core relationships include:

- Person to roles, skills, evidence, organisations, learning, opportunities and outcomes.
- Role to industry, seniority, skills, behaviours and compensation.
- Skill to roles, learning, evidence and adjacent skills.
- WRI dimension to evidence types, scenarios and role relevance.
- Organisation to roles, opportunities and hiring outcomes.
- Opportunity to role, skills, location, capacity, compensation and evidence generated.
- Learning to skills, roles, WRI gaps and outcomes.
- Evidence to source, verifier, dimensions, recency and confidence.
- Outcome to hire, completion, rehire, retention, performance and promotion.

GRAPH-001: Taxonomy changes are versioned.

GRAPH-002: Matching explanations can trace major graph signals.

GRAPH-003: AI can propose mappings, but authoritative ontology changes require validation.

25. Workforce Intelligence and Institutional Data

Workforce Intelligence is aggregate and privacy-protected by default.

Potential insight families include:

- Readiness bands by sector, geography or career family where cohort thresholds permit.
- Skills and capability gaps.
- Development patterns from Levav 28 and Learn.
- Time-to-opportunity and transition patterns.
- Employer demand.
- Work mode patterns.
- Volunteer and contribution patterns.
- QuickWork capacity, demand, completion and repeat engagement.
- Compensation coverage where valid data exists.

DATA-001: Institutional products do not provide default access to individual WRI records.

DATA-002: Minimum cohort thresholds suppress re-identifiable groups.

DATA-003: Exports are purpose-bound, time-bound and audited.

DATA-004: Every institutional dataset has governance metadata including:

- Owner.
- Purpose.
- Lawful basis or contract.
- Fields.
- Cohort rules.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Retention.
- Expiry.
- Access scope.

26. Monetisation Architecture

Levav's revenue model must reinforce product value without corrupting evidence or matching.

Potential revenue lines include:

- Employer subscriptions.
- Enterprise employer plans.
- QuickWork transaction or service fees.
- Talent Premium.
- Levav Learn.
- Sponsored development cohorts.
- Institutional analytics.
- Future scoped enterprise integrations.

MON-001: Paid talent features cannot override role relevance.

MON-002: Employer plans control entitlement, not WRI scoring.

MON-003: Institutional payment never buys private individual behavioural data.

MON-004: Commercial features must expose clear value. Do not add paywalls that block the user's core ability to build readiness evidence.

27. Trust, Fraud and Evidence Integrity

The system must assume that evidence can be manipulated and design for trust without becoming hostile to users.

Trust controls may include:

- Identity verification.
- Duplicate-account detection.
- Employer and organisation verification.
- Suspicious reciprocal review detection.
- Collusion signals.
- Fake QuickWork assignment detection.
- Reviewer trust scoring.
- Portfolio plagiarism or duplicated evidence detection where feasible.
- Uploaded evidence integrity checks.
- Payment risk controls.
- Manual review queue.

TRUST-001: Reviewer trust influences evidence confidence rather than becoming a public popularity score.

Reviewer trust may consider:

- Verified relationship context.
- History of completed reviews.
- Consistency with objective assignment data.
- Dispute history.
- Suspicious reciprocal behaviour.

TRUST-002: Disputes are first-class workflow states.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Evidence can be quarantined.
- Resolution is logged.
- Recalculation occurs through controlled events.
- User appeal does not erase audit history.

28. AI Architecture

AI is a product capability layer, not the source of truth for Levav's evidence system.

AI-001: Provider abstraction.

The system should avoid hard-wiring critical product logic to one model provider where a reasonable abstraction is possible.

AI-002: Structured outputs.

Critical AI functions should return validated structured outputs rather than uncontrolled prose when the result feeds product logic.

AI-003: Version prompts and evaluation logic.

Record relevant versions for:

- Scenario generation.
- Persona instructions.
- Response evaluation.
- Employer search parsing.
- Job design analysis.
- Interview guidance.
- Recommendation reasoning.

AI-004: Human confirmation where authority matters.

User or employer confirmation is required before AI inference becomes authoritative for profile, job requirements or other declared facts.

AI-005: Evaluation and regression testing.

Maintain test sets for important AI capabilities. New prompt or model versions should be evaluated against expected behaviour before broad release.

AI-006: Model monitoring.

Track:

- Latency.
- Failure rate.
- Structured output validity.
- Cost.
- Scenario failure patterns.
- Evaluation disagreement where human review exists.
- Drift in WRI-related evaluation.

AI-007: Graceful fallback.

AI failure must not corrupt evidence or silently produce fabricated facts.

AI-008: Guardrails by capability.

CV extraction:

- User confirmation.

Career profiling:

- Explain inference and allow correction.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

Scenario generation:

- Constrained by role, industry, level, competency and scenario schema.

Workplace personas:

- Role-bound context and isolated journey state.

Response evaluation:

- Rubric-based, versioned and confidence-aware.

Employer search parsing:

- Convert natural language into structured filters for confirmation.

Salary assistant:

- Interpret valid data, never invent benchmarks.

Recommendations:

- Use structured Workforce Graph context and explain major signals.

Integrity detection:

- Risk signal only.

Institutional summaries:

- Use authorised aggregate datasets only.

29. Security, Privacy and User Trust

SEC-001: Encrypt data in transit and at rest through the selected managed infrastructure.

SEC-002: Secrets stay server-side in managed secret storage.

SEC-003: Organisation data uses tenant-aware authorisation.

SEC-004: Protected readiness data uses server-side field and entitlement controls.

SEC-005: Log protected WRI views, exports, privileged admin changes and institutional access.

SEC-006: Privileged administration requires stronger authentication, including MFA where supported.

SEC-007: Uploaded evidence is scanned and file types are restricted.

SEC-008: Sensitive endpoints are rate-limited.

SEC-009: Raw assessment telemetry retention is separated from durable evidence retention.

SEC-010: Support lawful correction, export and deletion workflows.

SEC-011: Hidden scoring rubrics and internal weights are never exposed to employers.

SEC-012: Identifiable behavioural data is not sold as an institutional data product.

29.1 Privacy and Evidence Centre

PRIV-001: Every user should have a Privacy and Evidence Centre.

It should show:

- What evidence Levav holds.
- What is self-declared and what is verified.
- Who verified evidence.
- Which evidence can affect WRI.
- Public, employer-visible and private visibility.
- Protected WRI access history where policy permits.
- Dispute options.
- Consent and revocation controls where applicable.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

30. African Infrastructure Requirements

AFR-001: Mobile-first responsive design.

AFR-002: Low-data mode for feed and profile media.

AFR-003: Image compression and resumable uploads.

AFR-004: Graceful poor-network states.

AFR-005: Offline draft capability for safe non-proctored Levav 28 tasks.

AFR-006: Email or SMS fallback for high-value notifications where supported and justified.

AFR-007: Local currency, timezone and locale support.

AFR-008: Architecture ready for additional African languages.

AFR-009: Device and network limitations must not be mistaken for low readiness.

AFR-010: Accessibility is part of the product quality bar.

31. Current Technology Direction and Migration Rules

The current repository is the implementation source of truth for stack choices.

As of this PRD edition, the active direction includes:

- React 19.
- Vite 6.
- TypeScript 5.9.
- React Router.
- TanStack Query.
- Tailwind CSS.
- shadcn/ui and Radix primitives.
- framer-motion.
- Hono backend services.
- tRPC.
- Drizzle ORM.
- PostgreSQL and Supabase infrastructure.
- Supabase-linked identity and row-level security foundations in the current data model.
- Vitest automated tests.
- Vercel deployment.

ENG-001: Do not migrate to Next.js, Prisma or a second UI framework merely because an older design example used it.

ENG-002: Preserve the working authentication and organisation foundations.

ENG-003: Before changing a major subsystem, Codex verifies the current code and runtime rather than trusting stale documentation.

ENG-004: The product specification defines behaviour. Folder examples do not override the real repository structure.

ENG-005: Mock and localStorage surfaces must be connected incrementally to real services. Do not perform a blind rewrite of the entire app.

32. Backend Service Boundaries

The following are logical service responsibilities. Codex may implement them within the existing architecture without forcing microservices.

- Identity: auth, sessions, verification, access and organisation membership.
- Profile: career profile, CV extraction, skills and visibility.
- Organisation: organisation pages, membership and verification.
- Levav 28: journeys, scenarios, personas, attempts and feedback.
- Evidence: evidence nodes, verification, disputes and provenance.
- WRI: dimensions, scoring, confidence, snapshots and model versions.
- Workforce Graph: taxonomy and relationships.
- Matching: talent-role and opportunity relevance with explainability.
- QuickWork: capacity, listings, matching, assignments, evidence and reviews.
- Jobs: jobs, applications and pipeline.
- Employer Intelligence: search profiles, pools, comparisons and Hire Confidence.
- Compensation: role mapping, data adapters, benchmarks and confidence.
- Impact: volunteer opportunities, placements and verification.
- Learn: courses, enrolment, entitlements and progress.
- Social: posts, follows, comments, saves and reports.
- Messaging: threads, requests, context and anti-spam.
- Billing: plans, subscriptions, invoices, entitlements and payment adapters.
- Data Platform: events, aggregates and institutional products.
- Trust: moderation, fraud, disputes, risk and audit.
- Notifications: in-app, email or push orchestration and preferences.

33. Core Data Model

The data model should support at least these conceptual entities as the relevant modules are implemented:

- User.
- ProfessionalProfile.
- Organisation.
- OrganisationMembership.
- CareerTaxonomy.
- Levav28Journey.
- Scenario.
- ScenarioAttempt.
- Evidence.
- WRIDimension.
- WRISnapshot.
- RoleReadinessSnapshot.
- CapacityProfile.
- QuickWorkListing.
- QuickWorkAssignment.
- WorkEvidenceAsset.
- PerformanceReview.
- EmploymentEvidence.
- EmployerSearchProfile.
- TalentPool.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Job.
- Application.
- VolunteerOpportunity.
- VolunteerPlacement.
- Course.
- Enrolment.
- Subscription.
- Entitlement.
- HireOutcome.
- CompensationBenchmark.
- AuditEvent.

DATA-MODEL-001: Database design must be derived from current schema and migration history rather than pasted from the old v3.3 Prisma example.

DATA-MODEL-002: Every evidence-producing entity must be able to emit a controlled evidence event.

DATA-MODEL-003: Protected data visibility is represented in the data model rather than relying only on client UI hiding.

34. Event-Driven Evidence Pipeline

Evidence-producing UI actions must not directly mutate WRI scores.

Events should trigger controlled consumers.

Examples include:

- profile.confirmed -> matching and recommendations.
- levav28.attempt.completed -> evidence, WRI and feedback.
- quickwork.assignment.started -> attendance and audit.
- quickwork.assignment.submitted -> evidence and client review.
- quickwork.assignment.completed -> evidence, WRI, billing and profile.
- quickwork.review.completed -> evidence, WRI and reputation.
- employment.evidence.verified -> evidence, WRI and profile.
- volunteer.placement.verified -> evidence, WRI and impact.
- course.assessment.completed -> eligible evidence and recommendations.
- employer.wri.viewed -> audit and privacy monitoring.
- hire.outcome.recorded -> Quality of Hire and model evaluation.
- evidence.disputed -> trust workflow and WRI recalculation queue.

EVENT-001: Consumers should be idempotent for finalisation and payment-sensitive workflows.

EVENT-002: Evidence and WRI changes remain traceable to source events.

35. API Principles

API-001: Protected endpoints require authenticated identity and server-side authorisation.

API-002: Organisation endpoints require membership and tenant checks.

API-003: Protected WRI endpoints require entitlement and visibility checks.

API-004: Registry, payment, compensation and AI providers use adapters where practical.

API-005: Enterprise or institutional credentials are scoped, revocable, rate-limited and audited.

API-006: Billing, evidence finalisation and payment-sensitive actions are idempotent.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

API-007: Externally consumed schemas are versioned.

API-008: Error responses must not leak internal scoring logic or sensitive tenant information.

36. Required Product States

Every major surface must be designed for real states, not only the happy path.

Talent Home:

- New user.
- Provisional WRI.
- Active Levav 28.
- Evidence growing.
- Opportunity-ready.
- Employed development.

WRI:

- No data.
- Provisional.
- Medium confidence.
- High confidence.
- Disputed evidence.
- Historical trajectory.

Levav 28:

- Day 1.
- Active day.
- Missed day.
- Feedback.
- Integrity follow-up.
- Day 15 threshold.
- Day 28 report.

QuickWork Talent:

- No capacity.
- Capacity active.
- Matched.
- Active assignment.
- Submission.
- Review.
- Completed.
- Dispute.

QuickWork Client:

- Draft.
- Verification.
- Applicants or matched pool.
- Matched.
- Active.
- Review due.
- Completed.
- Dispute.

Employer Talent View:

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Public.
- WRI locked.
- Subscribed.
- Provisional WRI.
- High-confidence WRI.
- Hidden evidence.
- Plan limit.

Employer Search:

- No criteria.
- AI-parsed criteria.
- Results.
- Insufficient results.
- Saved pool.

Compensation:

- No data.
- Broad benchmark.
- Confident benchmark.
- Comparison.
- Insufficient data.

Organisation:

- Unverified.
- Verified.
- Opportunity active.
- Verification due.

Institutional:

- No contract.
- Scoped access.
- Cohort suppression.
- Export requested.
- Export expired.

37. Observability and Model Monitoring

OBS-001: Structured logs use correlation identifiers.

OBS-002: Product metrics include authentication, payment, WRI recalculation, AI latency, scenario failures and matching quality.

OBS-003: WRI model versions have calibration and drift monitoring.

OBS-004: Alert on unusual protected WRI access and institutional export activity.

OBS-005: Manual overrides, verification changes and disputes are auditable.

OBS-006: High-risk scoring and data changes use feature flags or controlled rollout.

OBS-007: Critical user journeys have operational dashboards or equivalent monitoring before production scale.

38. Success Metrics

Levav must optimise verified progression, not vanity activity.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

38.1 Primary ecosystem metric

Primary metric: Readiness Activation Rate.

Definition:

The percentage of eligible active users who move from a provisional state to an evidence-sufficient opportunity-ready state within a defined window, with sufficient diversity of E2 and E3 evidence for the relevant pathway.

This avoids rewarding the system for generating many trivial evidence nodes.

38.2 Supporting metrics

Talent:

- Onboarding completion.
- Levav 28 Day 7, Day 15 and Day 28 retention.
- Evidence growth.
- Time-to-first verified opportunity.
- Opportunity conversion.

Readiness:

- Confidence growth.
- Dimension improvement.
- Evidence diversity.
- E2 to E3 calibration.
- E3 to E4 validation.

Employed Talent:

- Current-role readiness engagement.
- Next-role progress.
- Verified employment evidence.
- Career progression outcomes where available.

QuickWork:

- Capacity listed.
- Fill rate.
- Completion rate.
- On-time rate.
- Rehire rate.
- Verified hours.
- Earnings.
- Conversion to recurring or larger opportunity.

Employer:

- Free-to-paid conversion.
- Time-to-shortlist.
- Search-to-contact.
- Time-to-hire.
- Employer retention.

Hiring Quality:

- Probation success.
- Retention.
- Time-to-productivity.
- Prediction calibration.

Compensation:

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Benchmark usage.
- Data coverage.
- Confidence distribution.

Learn:

- Enrolment.
- Completion.
- Evidence-backed improvement where supported.

Impact:

- Verified placements.
- Outputs.
- Completion.
- Contribution evidence.

Institutional:

- Data coverage.
- Freshness.
- Cohort validity.
- Report usage.
- Privacy incidents, target zero.

Trust:

- Disputes.
- Fraud rate.
- Moderation response.
- Unauthorised access attempts.

38.3 Secondary velocity metric

Readiness Velocity may be reported as verified evidence progression divided by time to first verified opportunity, but it must use weighted and quality-controlled evidence. Raw node count alone must never be a target.

39. Product Differentiation and Moat

Levav's advantage should come from the connected system rather than one isolated feature.

The product moat is the closed loop:

- Develop in realistic simulation.
- Create E2 evidence.
- Interpret evidence through versioned WRI.
- Convert readiness into QuickWork or other opportunity.
- Create E3 verified real-world evidence.
- Observe post-hire and longitudinal outcomes as E4.
- Use outcomes to improve matching, scenarios and WRI calibration.
- Aggregate privacy-protected patterns into Workforce Intelligence.

Important design lessons adopted from current industry practice include:

- Use job-relevant simulations instead of abstract quizzes alone.
- Use dynamic persona interactions and practical work artefacts.
- Use structured skill and role context for matching.
- Explain major matching signals.
- Help employers structure work and roles before matching.
- Support repeat practice and targeted development.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Connect flexible work to stronger career evidence.

Levav must combine these lessons with its own evidence provenance, longitudinal readiness loop, Africa-first operating conditions and capacity activation model.

No competitor brand, proprietary workflow or copied language belongs in the product specification or user experience.

Levav language must remain internally consistent even when product lessons are informed by external patterns.

40. Dual-Agent Operating Model

The project uses two primary agents with deliberately different authority.

Claude and Codex are collaborators. They are not two engineers editing the same system.

40.1 Claude role

Claude is Product Command.

Claude acts as:

- Product Manager.
- Programme Manager.
- Requirements owner.
- Product architecture reviewer.
- Domain logic spec owner.
- WRI scientific specification coordinator.
- Levav 28 scenario and rubric specification owner.
- Product language and UX writing system owner.
- Levav Impact product specification owner.
- Professional Feed and Network product specification owner.
- Acceptance-test designer.
- QA gatekeeper.
- Decision log owner.
- Cross-workstream coordinator.

Claude may use subagents for research, scenario review, psychometric review, security review, accessibility review, product QA and test design.

All Claude subagents inherit Claude's write boundaries.

Claude is primarily read-only over production code.

Claude must inspect implementation enough to review it, but must not become a competing coding agent.

40.2 Codex role

Codex is Engineering Command.

Codex acts as:

- Lead implementation engineer.
- Runtime owner.
- Frontend engineer.
- Product language implementation and copy-system engineer.
- Feed, profile and social interaction implementation engineer.
- Levav Impact implementation engineer.
- Backend engineer.
- Database and migration engineer.
- API engineer.
- Test engineer.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- CI and build engineer.
- Performance engineer.
- Security implementation engineer.
- Production debugging owner.

Codex owns executable implementation.

Codex must continue fixing a scoped work packet until the acceptance tests it controls pass or a genuine product decision is required.

40.3 Why this split exists

The earlier v3.3 split assigned Claude executable domain code and Codex runtime code. That creates unnecessary merge boundaries inside one feature.

The v4.1 split removes that risk.

One agent owns product meaning. One agent owns executable implementation.

41. Non-Overlap File Ownership

The repository should adopt explicit write ownership.

Claude exclusive write scope

Recommended paths:

- CLAUDE.md.
- docs/prd/.
- docs/product/.
- docs/decisions/.
- docs/qa/.
- docs/agent-handoffs/claude/.
- specs/scenarios/.
- specs/rubrics/.
- specs/product-contracts/.

Claude may read all repository files.

Claude does not edit:

- src/.
- server/.
- api/.
- db/.
- executable tests.
- build configuration.
- deployment configuration.

Codex exclusive write scope

Recommended paths:

- AGENTS.md.
- src/.
- server/.
- api/.
- db/.
- tests/ or existing test locations.
- scripts/.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- package.json and lockfile where required.
- TypeScript, Vite and build configuration.
- CI configuration.
- deployment configuration.
- docs/implementation/.
- docs/agent-handoffs/codex/.

Codex may read all product documents and specs.

Codex does not edit:

- Master PRD.
- Product acceptance criteria.
- WRI scientific definitions.
- Approved scenario rubrics.
- Product Decision Records.

41.1 Shared control files

A file must never have simultaneous write ownership.

If content must move from Claude-owned specification into a Codex-owned root instruction file, Claude records the approved text in a handoff. Codex applies it.

42. Agent Handoff Protocol

Every implementation unit begins with a Claude Work Packet.

42.1 Claude Work Packet format

Each packet contains:

- Work Packet ID.
- Related PRD requirement IDs.
- Product problem.
- User journey.
- In scope.
- Out of scope.
- Existing behaviour to preserve.
- Functional acceptance criteria.
- Data and privacy requirements.
- Expected events.
- Required UI states.
- Test scenarios.
- Dependencies.
- Open product decisions, if any.

Status begins as READY_FOR_BUILD.

42.2 Codex implementation report

Codex responds with:

- Work Packet ID.
- Audit classification of affected existing code.
- Implementation summary.
- Files changed.
- Database migration summary.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- API changes.
- Tests added or changed.
- Commands run and results.
- Security and permission checks.
- Performance or low-bandwidth considerations.
- Known limitations.
- Status: `READY_FOR_REVIEW` or `BLOCKED_PRODUCT_DECISION`.

42.3 Claude acceptance review

Claude checks:

- Requirement coverage.
- Acceptance criteria.
- Product semantics.
- Edge states.
- Evidence and WRI integrity.
- Privacy and entitlement behaviour.
- User experience.
- Test coverage.
- Regression risk.

Claude returns either:

- `ACCEPTED`.
- `DEFECTS_FOUND` with a numbered defect list.
- `BLOCKED_PRODUCT_DECISION` with a decision record.

Codex fixes `DEFECTS_FOUND` and returns the same Work Packet to `READY_FOR_REVIEW`.

The cycle repeats until `ACCEPTED`.

43. Parallel Work Rules

Parallelism is allowed only when workstreams are technically disjoint or their contracts are stable.

Claude controls work packet sequencing.

The following are prohibited:

- Claude and Codex editing the same production file.
- Two Codex branches modifying the same database migration chain without coordination.
- Building WRI before the Evidence Graph contract is stable.
- Building premium WRI UI without server-side entitlement design.
- Building Workforce Intelligence on raw private individual data without governance controls.

A dependency must be resolved before dependent work is accepted.

44. Claude Subagent Model

Claude may create specialised subagents to increase quality and speed.

Recommended roles:

Product research reviewer

Checks current patterns, user needs and product gaps. Produces research notes only.

Levav 28 scenario reviewer

Checks realism, difficulty, role relevance, construct coverage and ambiguity.

WRI scientific reviewer

Checks behavioural anchors, construct definitions, scoring proposals, confidence and validation design.

Trust and privacy reviewer

Checks consent, evidence visibility, disputes, institutional aggregation and employer access.

Accessibility and low-bandwidth reviewer

Checks mobile, network degradation, media dependence, accessibility and device assumptions.

QA acceptance reviewer

Turns acceptance criteria into test scenarios and checks Codex implementation reports.

Subagents do not change production source files.

45. Codex Engineering Behaviour

CODEX-001: Audit before edit.

Codex reads the current implementation and classifies affected parts before coding.

CODEX-002: Preserve working authentication.

Do not weaken or replace auth to make a feature easier.

CODEX-003: Implement vertical slices.

Prefer a functioning end-to-end slice over broad disconnected scaffolding.

CODEX-004: Test business logic.

New WRI, matching, payment, entitlement, evidence, transition and security logic requires automated tests.

CODEX-005: Fix until green.

For a work packet, Codex runs the relevant tests, type checks and build. It fixes regressions within scope before declaring READY_FOR_REVIEW.

CODEX-006: No product invention by code.

If the PRD leaves a product decision unresolved, Codex chooses the safest reversible implementation where possible and records the decision needed. It does not silently invent WRI coefficients, paywall rules, compensation data or evidence policy.

CODEX-007: Keep implementation truth current.

After accepted work, update docs/implementation/IMPLEMENTATION_STATE.md or the approved equivalent with factual current state.

46. Build Sequence

The build should prove the closed readiness loop early.

Sprint 0: Control Plane and Repository Truth

Claude:

- Publish Master PRD v4.1 in the repository.
- Replace stale Claude instructions with the v4 role.
- Create requirements index.

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Create decision log structure.
- Create QA and handoff templates.

Codex:

- Audit current main branch.
- Update AGENTS.md to point to the v4 authority.
- Verify real auth, organisation, deployment, database and test state.
- Repair stale implementation documentation.
- Ensure typecheck, test and build scripts reflect actual commands.
- Add lint as a gate only after it is correctly configured.
- Establish CI if missing.

Exit gate:

- Both agents agree on current implementation truth.
- No stale document can override the Master PRD.
- Auth remains working.
- Typecheck, tests and build have known verified states.

Sprint 1: Identity, Taxonomy and Personalisation

Goal:

- Complete the person-owned identity model around the existing auth foundation.
- Strengthen organisation membership and verification.
- Implement career taxonomy and intelligent onboarding.
- Connect Personal Home and professional profile to real services.

Exit gate:

A new user can register, confirm identity, complete intelligent onboarding, receive a real personal profile state and enter the relevant next journey without mock identity logic.

Sprint 2: Evidence Graph Foundation

Goal:

- Build evidence entities, provenance, visibility, disputes and controlled event ingestion.
- Establish E0-E4 types.
- Separate factual observations from subjective evaluations.
- Build Privacy and Evidence Centre foundations.

Exit gate:

A verified event can create an auditable evidence node and a dispute can change its processing state without deleting history.

Sprint 3: Versioned WRI Engine

Goal:

- Build WRI dimension definitions.
- Configurable scoring.
- Evidence confidence.
- Evidence coverage.
- Trajectory.
- Snapshot versioning.
- Role Readiness separation.
- Auditability.

Exit gate:

Given controlled test evidence, the system produces deterministic versioned WRI snapshots and separate confidence values, with traceable contributions and tests.

Sprint 4: Levav 28 Vertical Slice

First target:

- Days 1-5 end to end.
- Persistent journey state.
- At least two scenario modalities.
- Persona continuity.
- Rubric-based evaluation.
- E2 evidence creation.
- Feedback and coaching split.
- Poor-network state.

Then expand to Days 6-28 after the vertical slice is accepted.

Exit gate:

A real user can complete a controlled Levav 28 task and create valid E2 evidence that updates WRI through the event pipeline.

Sprint 5: QuickWork Closed Loop

Build:

- Capacity Profile.
- Client work scoping.
- Listing.
- Matching.
- Assignment lifecycle.
- Submission.
- Client review.
- Work evidence.
- Dispute state.
- Client reputation.
- Payment state model.

Exit gate:

A startup can define work, match to eligible capacity, complete the assignment, close it out and create E3 evidence without manual database manipulation.

Sprint 6: Employer Intelligence

Build:

- Employer verification completion.
- Talent Search Profile.
- Natural-language criteria parsing and confirmation.
- Explainable candidate search.
- WRI entitlement.
- Candidate comparison.
- Saved pools.
- Job Design Assistant.
- Interview Intelligence.
- Hiring pipeline.

Exit gate:

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

A verified subscribed employer can discover a relevant candidate, understand why the candidate appears, view only entitled readiness data and progress the candidate through a real pipeline.

Sprint 7: Post-Hire Outcome and E4

Build:

- Probation outcome capture.
- Retention and time-to-productivity fields.
- Quality of Hire.
- Prediction versus outcome comparison.
- E4 evidence ingestion.

Exit gate:

A real employer-confirmed outcome can create E4 evidence and enter model evaluation without rewriting historical WRI snapshots.

Sprint 8: Learn, Impact, Champions and Feed Integration

The goal is not to build isolated content islands. Sprint 8 makes Levav feel like a living professional ecosystem while preserving evidence integrity.

Build or complete:

- Levav Learn real services and evidence-safe achievement flows.
- Levav Impact opportunities, placements, verification, contribution records and organisation workspace.
- Champions verification and contribution surfaces.
- Professional profile editing, cover/background image and featured content.
- Secure email-change account flow that preserves one Levav identity.
- Feed publishing, organisation posts, media, reactions, comments, follows, saves, reporting, notifications and discovery views.
- What is New and sourced workforce/news cards without fabricated AI news.
- Consistent Levav language and copy architecture across the surfaces implemented to this point.

Each surface must connect into:

- Identity.
- Evidence.
- Recommendations.
- Workforce Graph.
- Visibility policy.
- Moderation and trust.

Exit gate:

A member can maintain a rich professional presence, publish and engage in a real feed, participate in verified Impact, receive useful discovery, and see consistent Levav language across the system. Only eligible verified evidence influences WRI.

Sprint 9: Workforce Graph, Compensation and Workforce Intelligence

Build graph services and stronger matching first.

Activate Compensation Intelligence only for supported data.

Activate aggregate Workforce Intelligence only after governance and cohort rules are operational.

Exit gate:

Institutional insight cannot expose private individual data, and compensation never fabricates a benchmark.

Sprint 10: Production Hardening

Complete:

Develop readiness. Prove readiness. Convert readiness into opportunity. Learn from outcomes.

- Admin workflows.
- Moderation.
- Fraud controls.
- Disputes.
- Audit.
- Observability.
- Performance.
- Route-level code splitting.
- Mobile QA.
- Low-bandwidth QA.
- Accessibility.
- Security hardening.
- Backup and recovery procedures where relevant.

47. Definition of Done

A feature is not done because it renders.

A Work Packet is done only when:

- Acceptance criteria are met.
- Existing valid functionality is not regressed.
- Server-side permission and entitlement checks are tested.
- Database migration is documented and reversible where practical.
- Loading, empty, error and permission states exist.
- Mobile behaviour is verified.
- Poor-network behaviour is considered where relevant.
- Critical business logic has automated tests.
- Audit and analytics events exist where required.
- Security-sensitive paths have negative tests.
- Admin or support workflow is documented where required.
- Product copy does not overclaim scientific certainty.
- AI output failure states are handled.
- Evidence changes are auditable.
- Protected WRI cannot be bypassed through direct URL or API calls.
- Claude has marked the Work Packet ACCEPTED.

47.1 Current engineering commands

Codex must use the actual repository commands, not obsolete PRD examples.

Current expected gates include:

- npm run typecheck
- npm run test
- npm run build

Lint becomes a required gate once a working lint script and dependencies are intentionally configured.

48. Product and Engineering Never-Do List

The following actions are prohibited without an explicit approved product decision:

- Rebuild working authentication from scratch.

- Create duplicate person identities for different workspaces.
- Publish protected WRI publicly by default.
- Conflate WRI with Role Fit.
- Hide evidence confidence.
- Hard-code provisional WRI coefficients as permanent scientific truth.
- Let UI requests directly mutate WRI scores.
- Penalise a user solely because an AI detector flagged a response.
- Use protected characteristics in employment ranking.
- Treat social engagement as readiness.
- Let paid visibility override relevance.
- Fabricate compensation benchmarks.
- Sell identifiable behavioural data to institutions.
- Turn QuickWork into an unstructured gig board.
- Turn volunteering into disguised unpaid employment.
- Add invasive workforce surveillance as the default.
- Introduce a second UI framework without clear migration approval.
- Migrate stack merely to satisfy an old architecture example.
- Build large disconnected feature scaffolds before proving vertical slices.
- Allow Claude and Codex to edit the same files concurrently.

49. Launch Decisions Still Requiring Human Approval

The following remain configuration or policy decisions. Agents may prepare options but should not silently decide them as permanent policy:

- Final WRI scale.
- Final evidence weights.
- Confidence thresholds.
- Recency rules.
- Scientific validation thresholds.
- Day 15 Levav 28 evidence-sufficiency threshold.
- Final scenario libraries.
- QuickWork pricing.
- Platform transaction fee.
- Payment and escrow provider model.
- Dispute rules.
- Employer plan names and prices.
- Employer seats and WRI depth.
- Talent Premium price and benefits.
- Launch payment providers.
- Organisation registry sources.
- Compensation data providers.
- Supported compensation countries and cities.
- Institutional cohort thresholds.
- Institutional pricing and licensing.
- Champion content rights and revenue sharing.

50. Agent Start Instructions

50.1 Claude first instruction

Read this Master PRD completely before issuing work.

Act as Levav Product Command.

Your first job is not to code.

Your first job is to:

1. Audit the current repository and current implementation reports against this PRD.
2. Create a requirement coverage matrix using KEEP, ENHANCE, MODIFY, COMPLETE, BUILD, REMOVE and DEFER.
3. Identify stale documentation that could mislead either agent.
4. Create the Sprint 0 Work Packets for Codex.
5. Establish the product decision log.
6. Establish the acceptance and QA structure.
7. Review Codex's completed work packets against this PRD.
8. Use subagents for research and specialist review where useful.
9. Do not edit production code.
10. Do not mark work accepted until evidence supports the claim.

50.2 Codex first instruction

Read this Master PRD completely before implementation.

Act as Levav Engineering Command.

Your first job is to:

1. Inspect main and verify the current runtime before trusting old docs.
2. Preserve the working authentication and organisation foundations.
3. Audit current files against the first Claude Work Packet.
4. Implement only the approved scope.
5. Own all executable code, migrations, tests, build and deployment changes.
6. Run typecheck, tests and build before READY_FOR_REVIEW.
7. Fix scoped defects until the work packet passes.
8. Update factual implementation-state documentation after accepted work.
9. Do not change product rules or scientific methodology silently.
10. Never edit Claude-owned product specification files.

51. Final Product Thesis

Levav develops, measures and verifies workforce readiness.

For talent, it turns development and productive activity into visible, verified evidence.

For employed professionals, it measures current-role readiness, next-role readiness and development trajectory.

For startups, QuickWork turns flexible capacity into affordable execution and verified real-world work evidence.

For employers, Levav reduces hiring uncertainty by showing the person beyond the CV.

For Africa's workforce ecosystem, Levav turns privacy-protected evidence into intelligence about readiness, gaps, opportunity, compensation and development.

For contributors and mission-driven organisations, Levav Impact turns verified service and real outputs into visible contribution without confusing volunteering with employment.

For the professional community, the Feed makes Levav a place people can return to daily to publish work, discover opportunities, learn, follow organisations, see trusted updates and build professional identity.

Across every surface, Levav must use one clear product language so the experience feels coherent from the first landing-page sentence to the deepest workforce-intelligence workflow.

The product should be judged by one question:

Does Levav help a person become more ready, prove that readiness, convert it into meaningful opportunity, and help the workforce system learn from the outcome?